

SinhalaJournal-LLM: Data-Centric Adaptation and Reliable Evaluation of a Sinhala Language Model for Journalism

Abstract—Sinhala has limited Natural Language Processing support for practical newsroom writing. We adapt one Sinhala language model for grammar correction, headline generation, news summarization, and journalistic style rewriting. We collected approximately 884,193 Sinhala news articles from ten public Sri Lankan news websites between 2008 and March 2026 using crawlers developed for this project. After exact duplicate removal and quality filtering, 665,887 articles remained. We trained a separate LoRA adapter for each task on the same SinLlama base. On the 286 input grammar benchmark, the final adapter corrected 44.06% of sentences that needed a change and preserved 69.93% of already correct controls. We replayed the current safety validator on Stage 6 and found that it kept the neural result unchanged. In headline generation, duplication before data splitting increased ROUGE 1 from 0.1382 to 0.2933. During our summarization audit, we found about 81% overlap between a later test set and legacy training through other length versions of the same articles. We then retrained both recipes with article level separation and obtained very similar results on 273 held out articles. In our frozen style test of 500 outputs, we observed 92.89% number preservation and 66.19% F1 from a simple style classifier. We also conducted a field study with 13 participants. They rated meaning preservation, natural writing, and usefulness at 4.15, 4.23, and 4.38 out of 5. The results show that target quality, correct data splits, prompt consistency, and careful evaluation are important when adapting language models for low resource Sinhala generation.

Index Terms—Sinhala NLP, low resource NLP, natural language generation, large language models, LoRA

I. INTRODUCTION

Large language models have improved text generation, but Sinhala still has fewer datasets, benchmarks, and practical tools than high resource languages [1], [2]. Sinhala specific encoders have improved classification [3], [4]. SinLlama extends Llama 3 with a Sinhala vocabulary and further Sinhala training [5], [6]. NSina also provides a large Sinhala news corpus and a headline benchmark [7]. Earlier Sinhala studies have covered tasks such as extractive summarization [8] and grammar correction [9]. However, a newsroom also needs one practical workflow for editing, headline writing, summarizing, and rewriting text while keeping names, numbers, dates, quotations, and meaning correct.

We use SinLlama as the shared base and train a separate LoRA adapter [10] for each task. We load the base in 4 bit form using the main idea of QLoRA [11]. Our main focus is not only the final scores. We study how target quality, data splitting, prompts, and evaluation choices can change the apparent model performance.

We address three research questions. First, can one shared Sinhala base model learn four different newsroom tasks through separate adapters? Second, how much can target quality and data splitting change the measured score? Third, what checks are needed before a higher score can be treated as a real improvement? These questions focus the study on reliable Sinhala generation.

Our work makes four main contributions. (1) We build and audit a large Sinhala news data pipeline and create task data for four newsroom tasks. (2) We adapt one Sinhala decoder to grammar correction, headline generation, summarization, and style rewriting. (3) We study data problems that changed measured results, including headline duplication and summarization split leakage. (4) We report task specific evaluation with a frozen grammar comparison, a Stage 6 safety replay, article level summarization tests, a frozen automatic style test, and a small human field study.

II. RELATED WORK

LoRA trains small low rank updates while the main model stays frozen [10]. QLoRA makes this more practical by combining adapters with a quantized base model [11]. These methods are useful when several tasks need to share one large model. SinLlama provides a Sinhala adapted decoder that is suitable for this setting [6].

Sinhala generation research is still limited. NSina provides data and headline experiments for Sinhala news [7]. Earlier Sinhala summarization work tested TF IDF and TextRank methods [8]. Multilingual encoder decoder models such as mT5 and XL Sum are also important references for abstractive generation [12], [13]. ByT5 provides a byte based alternative that can avoid some vocabulary problems [14]. For evaluation, ROUGE measures overlap [15]. BERTScore gives a semantic similarity signal using contextual representations [16]. Style rewriting needs both style change and content preservation, so automatic overlap alone is not enough [17].

Our datasets and prompts are different from the benchmarks in earlier Sinhala work. We therefore use this related work to position our tasks and model choices. We do not treat differences in reported scores across different datasets as a direct state of the art comparison.

TABLE I
TASK DATA AND MAIN EVALUATION EVIDENCE

Task	Final data	Main evaluation evidence
Grammar	36,006 rows	286 inputs with faulty and correct pairs
Headline	43,737 train	leakage test with leaked and repaired splits
Summary	35,547 articles	273 complete test articles and 819 outputs per recipe
Style	7,555 rewrites	100 unseen articles in five styles and 500 outputs

III. DATA AND METHOD

A. Corpus and Task Data

We collected approximately 884,193 Sinhala news articles from publicly available Sri Lankan news pages between 13 March 2008 and 4 March 2026 using crawlers developed for this project. The sources were Siyatha, LankaTruth, Hiru News, Vikalpa, Ada Derana, ITN News, Divaina, Dinamina, Sinhala News LK, and Vidusara. We did not use NSina articles to build these task datasets. We stored the category, date, headline, article content, source, and URL for the collected records. We removed exact duplicates and applied quality filters. This removed about 218,300 records and left 665,887 articles. We accepted articles with a body length from 300 to 10,000 characters, a title length from 10 to 200 characters, and at least 30% Sinhala characters. We removed rows that failed these checks instead of cutting them to fit.

The main news categories used in the project include Business, Politics, International, Sports, Entertainment, Education, and Local news. The category sizes are not equal. This is one reason the headline work tested category balancing. It also explains why balancing must be done only after the train and validation split.

Sinhala text cleaning required care. Soft hyphen U+00AD is invisible and can break tokenization, so it was removed. Zero width joiner U+200D was kept because it appears inside valid Sinhala letter forms. We also found punctuation and encoding problems that changed sentence boundaries. This showed that general removal of invisible characters is not safe for Sinhala.

Table I shows the final task data. Grammar includes faulty examples and correct controls. Summarization stores short, medium, and long targets for each article. We rebuilt the style dataset after early generated targets copied source text and repeated the same templates. Only 7,555 of 22,236 later style candidates passed the final checks.

B. Data Quality Changes

Each task had a different data problem. The final grammar training file contains 36,006 inputs. It covers spelling and orthographic errors, verb and agreement errors, word order, spacing, punctuation, register, pronouns, negation, numerals, and already correct controls. During the last target repair step, 4,999 target outputs were changed while the input count stayed the same. This gave a useful way to study target quality without adding more training examples.

The first headline balancing method copied minority category rows before the train and validation split. This allowed the same article and headline pair to appear on both sides. The corrected method splits real articles first and only balances the training side.

The raw summarization data contain 35,569 articles with three target lengths. During our data audit, we found rare joined word errors and number unit errors. We removed twenty two affected rows and kept 35,547 articles in the cleaned data. We used this small cleaning step to remove high impact factual problems rather than to increase ROUGE.

Style target generation had a larger quality problem. Early targets copied the input or reused fixed phrases. During our first audit, we found about 15% to 17% near verbatim copying in formal news and sports targets. We also checked source containment and found that most targets in that early file included the source article text directly. The adapter trained on those targets repeated the same behaviour, so simply generating more targets did not solve it. We then generated new targets and checked source copying, repeated phrases, Sinhala script quality, length, structure, numbers, and important facts. Only 34.0% of the 22,236 candidates passed.

C. Shared Model Design

The system keeps one SinLlama base and loads one task adapter at a time. This keeps the main Sinhala language knowledge shared while each task can learn its own output rules. Fig. 1 shows the model design used by the four components.




Fig. 1. Shared model design used by the four newsroom tasks

D. Adapter Setup and Prompts

All adapters use the same seven attention and MLP target modules. Table II reports the selected configurations, including the substantially smaller grammar adapter.

The target modules are `q_proj`, `k_proj`, `v_proj`, `o_proj`, `gate_proj`, `up_proj`, and `down_proj`.

Grammar asks the model to make only needed corrections and return correct text unchanged. Headline prompts include the category and requested word range. Summarization asks for short, medium, or long output and tells the model to use only article information. Style prompts request one of five news styles and include rules to keep entities, numbers, dates, quotations, and attribution. We verified that the current style training prompt and main serving prompt are the same for all five styles. Our internal headline evaluator did not use the

TABLE II
SELECTED LORA CONFIGURATIONS. ALL USE A 4 BIT NF4 SINLLAMA BASE WITH BF16 COMPUTE AND THE SAME SEVEN TARGET MODULES.

Task	Training data	$r/\alpha/\text{dropout}$	Max seq.	Eff. batch	LR	Epochs	Trainable params
Grammar	36,006 rows	4/4/0.05	512	8	5×10^{-5}	3	10,485,760
Headline	43,737 rows	64/128/0.08	768	8	5×10^{-5}	early stop	167,772,160
Summary	28,455 train articles	32/64/0.05	2048	16	2×10^{-4}	3	83,886,080
Style	7,555 rewrites	24/48/0.05	4096	16	1×10^{-4}	3	62,914,560

exact production wrapper, so we do not use its unstable final score as a production result.

During our summarization tests, we found a Sinhala specific decoding problem. The common setting `no_repeat_ngram_size=3` sometimes blocked a valid opening phrase because the same phrase was already present in the article prompt. In one of our stored outputs, this removed the first grapheme of the opening word. We therefore removed this setting and used a mild repetition penalty instead. This shows that a decoding rule that works for many languages can still damage Sinhala output when it acts across the prompt and generated text.

All adapter training used one fixed seed. This makes the saved training setup repeatable, but it does not measure training variation across different seeds. The final grammar setting changes both rank and LoRA scaling, so we describe it as a smaller parameter setup rather than a pure rank test.

We trained the final grammar adapter for about 6.36 hours on one NVIDIA A40. Each leakage controlled summarization retrain took about 40 hours on the same class of GPU. We do not estimate missing runtime records for the other final adapters.

E. Evaluation Protocol

Grammar Stage 6 has 286 inputs. There are 143 cases that need a correction and 143 paired controls that should stay unchanged. The faulty side contains 55 natural errors found in published news text, 80 injected unseen orthography cases, and eight spacing or punctuation cases. No normalized Stage 6 input appears in the final training data or Stage 5. Eighty seven changed cases contain exact correction pairs that are absent from training. The same inputs are used for the final adapter, a previous adapter, ByT5 small [14], mT5 small [12], and the unadapted SinLlama base. The ByT5 and mT5 baseline data use 28,519 training cases and 1,518 development cases, with 330 bridge cases removed to protect the rare correction pairs. We retained the three epoch ByT5 training record. We evaluated mT5 on Stage 6 using checkpoint 1783, although we did not retain its complete training log. We report overall exact match, change needed exact match, clean preservation, edit $F_{0.5}$, and recall on unseen correction pairs. Overall confidence intervals use 2,000 document clustered bootstrap samples. One member of our team manually checked the benchmark before this comparison. We did not complete an independent review with several raters.

We tested the grammar validator by replaying the same saved Stage 6 neural outputs through deterministic rule poli-

cies. This means the model output does not change between policy tests.

For headline generation, we focus on the verified data leakage experiment. When we reviewed our saved final results, we found two runs that differed by 23.78 percentage points in overall requested length accuracy. Our generation code also did not fully fix the random seed used for token sampling. Because of this, we do not report that unstable number as a final model accuracy result.

For summarization, we created a fixed 80/10/10 split at article level. It contains 28,455 training articles, 3,556 validation articles, and 3,558 test articles. A fixed set of 300 test articles was selected for repeated evaluation. Twenty seven rows did not have complete article content and all three reference summaries. They were removed by the existing completeness check. The final comparison therefore uses 273 articles and 819 generations per adapter. We report ROUGE 1, ROUGE 2, ROUGE L, requested length accuracy, clean sentence endings, joined word errors, and number unit errors. A smaller frozen set of 30 articles also uses XLM RoBERTa large BERTScore.

For style evaluation, we selected 100 unseen source articles after excluding the URLs used by the style training and target generation pipeline. We rewrote each article in five styles. This gave 500 adapter outputs and 500 no adapter outputs with the same prompts. We measured important fact coverage, all fact token coverage, number preservation, content similarity, and a simple character trigram style classifier. We do not call these values human style accuracy. We also conducted a field study with 13 users of the integrated system and collected five point ratings.

We define the task specific metrics here. Grammar clean preservation is the share of already correct controls that are returned unchanged. Edit $F_{0.5}$ gives more weight to precision because an unnecessary newsroom edit can be more harmful than leaving a small error. Summarization uses the generated word count divided by source word count. The short range is 0.04 to 0.18, the medium range is 0.12 to 0.32, and the long range is 0.22 to 0.55. A clean ending means that the output ends at a sentence boundary instead of being stopped by the token limit. The style fact and number measures are automatic preservation checks. The character trigram classifier is only a surface style signal.

IV. RESULTS AND DISCUSSION

A. Grammar Correction

Table III shows the Stage 6 model results. Table IV shows the safety replay. The unchanged input floor is 50% because

TABLE III
GRAMMAR STAGE 6 RESULTS ON 286 INPUTS

System	Overall [95% CI]	Change	Clean	$F_{0.5}$	Unseen
Final adapter	56.99 [50.00,63.99]	44.06	69.93	45.36	40.23
Previous adapter	57.34 [50.35,64.69]	44.06	70.63	46.21	41.38
ByT5 small	50.00	8.39	91.61	20.34	13.79
mT5 plus safety fallback	50.00	0.00	100.00	0.00	0.00
Unadapted SinLlama	0.00	0.00	0.00	0.00	0.00
Do nothing	50.00	0.00	100.00	0.00	0.00

TABLE IV
STAGE 6 GRAMMAR SAFETY REPLAY

Policy	Overall	Change	Clean
Neural output	56.99%	44.06%	69.93%
Rules only	50.00%	0.00%	100.00%
Older veto	57.34%	42.66%	72.03%
Current validator	56.99%	44.06%	69.93%

half of the inputs are already correct. The more useful grammar result is that the final adapter corrects 63 of 143 faulty inputs, which is 44.06%, while preserving 69.93% of the already correct controls. Its overall exact match is 56.99%. The previous adapter reaches 57.34%. Their overall 95% intervals overlap. One example is not enough to claim a performance difference. We retain the final adapter because it uses far fewer trainable parameters.

The unadapted base mainly failed to follow the task format. Seventy of 286 outputs kept response scaffold text and 217 outputs were more than twice the input length because of repetition. It also changed protected text in 12.94% of cases and changed numbers in 22.38% of cases. These failures show that task adaptation is needed for reliable correction. They do not prove that every language decision made by the adapter is better. ByT5 corrected 12 faulty cases but damaged 12 clean controls. All 286 raw mT5 candidates were rejected by its safety filter. Its reported 50% is therefore the unchanged input fallback and not raw mT5 accuracy.

In the safety replay, rules alone keep every clean control but do not correct any faulty case. The older veto policy changes 16 neural outputs. It increases clean preservation but lowers change needed exact match. The current recalibrated policy changes none of the 286 outputs, so it keeps the neural score unchanged. This means the current validator gives safety checks, but it does not add a measured strict accuracy gain on Stage 6.

A separate development comparison also shows why target quality mattered. Saved outputs from two earlier grammar adapters were scored using both the old target convention and the repaired one. Before the repair, the two runs scored 59.7% and 61.7% against the old gold. Against the repaired gold they scored 57.8% and 65.6%. The strictest comparison uses 143 cases where the gold text did not change and both systems received the same input. Exact match increased from 61.5% to 66.4% on this subset. This shows a broader correction gain as well as better agreement with the repaired target convention. These were successive development runs, so we do not claim that all of the change was caused only by target repair.

B. Headline Data Leakage

In an early headline experiment, we balanced minority categories by duplicating rows before splitting the data. The leaked evaluation reached ROUGE 1 of 0.2933 and 485 exact matches among 6,732 outputs. After we changed the pipeline to split real articles first, ROUGE 1 fell to 0.1382 and exact matches fell to zero among 4,795 outputs.

TABLE V
SELECTED HEADLINE CATEGORIES BEFORE AND AFTER SPLIT REPAIR

Category	Real n	Copies	Leaked R1	Repaired R1
Technology	10	40.0×	.7275	.1320
Editorial	86	4.7×	.6339	.0338
Health	48	8.3×	.6044	.1467
Science	75	5.3×	.4843	.1014
Sports	722	1.0×	.1713	.1646
Politics	722	1.0×	.1502	.1292

Table V shows the largest apparent gains in categories with small validation sets and heavy duplication. Sports and Politics did not need copying and changed much less. Rebuilding the split also changed which real articles appeared in validation, so not every small change can be blamed only on copying.

We do not use the later final headline score as a main result. Two saved runs disagree strongly, token sampling was not fully seeded, and our internal evaluation prompt did not exactly match the production prompt. We therefore report the leakage finding and defer the final model score to a controlled evaluation.

C. Summarization After Split Repair

The old summarization script first expanded each article into short, medium, and long examples and then used an 85/15 sample split. For an article with all three valid length examples, the chance that all three fall outside training is $0.15^3 = 0.003375$, which is about 0.34%. This calculation only applies to articles that contributed all three examples after filtering. The later audit measured 80.69% and 80.80% contamination over all articles in the frozen test set. The values use different groups of articles, so there is no conflict between the 0.34% calculation and the measured 81% overlap.

We retrained both summarization recipes using the same article level split. Table VI shows very similar results. ROUGE L differs by no more than 0.0018 across the three length ranges. Requested length accuracy stays between 86.8% and 94.1%. Clean endings are above 94% in every case. Both models have zero joined word errors on the full test.

TABLE VI
SUMMARIZATION ON 273 HELD OUT ARTICLES WITH 819 OUTPUTS PER RECIPE

Length	Recipe	R1	R2	RL	In band	Clean end
Short	raw	.7269	.4808	.5077	93.4%	98.5%
	clean	.7292	.4858	.5077	94.1%	98.9%
Medium	raw	.7558	.4804	.4577	92.7%	95.2%
	clean	.7492	.4812	.4559	90.5%	95.2%
Long	raw	.7805	.5208	.4566	86.8%	94.9%
	clean	.7805	.5265	.4575	88.6%	94.1%

A second frozen test on 30 articles adds BERTScore. The cleaned recipe is slightly higher for short, medium, and long outputs with F1 values of 0.9168, 0.9085, and 0.9053. The raw recipe scores 0.9118, 0.9071, and 0.9040. One short raw output contains a number scale error while the cleaned recipe has none on this small set. This supports the cleaning step as a useful safety measure, but the small difference does not prove that the cleaned recipe is generally better.

The automatic scores still miss some important factual errors. In one of our stored long summaries, a source reference to Portugal and Uruguay was replaced by Pakistan. During our teacher data audit, we also found number unit changes where the digits stayed the same but the scale word changed, such as a value expressed in millions being rewritten as lakhs. A fluent summary can therefore have good ROUGE and still change the actor, country, or size of an event. For this reason, we report number checks separately and keep stronger entity and event factuality testing as future work.

We exclude our earlier mT5 and TextRank summary results because we did not retain enough information about their sample size, exact test set, and evaluation setup to reproduce them reliably. We therefore avoid unsupported comparison claims. Entity and event factuality also need stronger testing because number checks do not catch every possible factual change.

D. Style Rewriting and Human Feedback

The final style system supports formal news, editorial, sports, youth, and feature writing. The frozen automatic test used 100 unseen articles and created 500 outputs from the adapter. The same prompts were also used with the no adapter base. Table VII shows a large gap in content and fact preservation. The adapter keeps 92.89% of numbers and 78.70% of important fact tokens. The simple style classifier reaches 66.19% F1 for the adapter and 22.00% for the no adapter base. The no adapter outputs also show strong repetition and weak task following. The large gap supports the benefit of task adaptation, but it does not establish final editorial quality. This classifier only measures surface style patterns, so we do not call it human style accuracy.

TABLE VII
FROZEN AUTOMATIC STYLE TEST ON 500 OUTPUTS

Measure	Adapter	No adapter
Important fact coverage	78.70	17.45
All fact token coverage	77.64	18.37
Number preservation	92.89	29.31
Content similarity	85.54	32.08
Style classifier F1	66.19%	22.00%

We also collected user feedback to complement the automatic style evaluation. Thirteen users tested the integrated system. They included four editors, five journalism students, two school teachers, and two other students. Mean style ratings were 4.15 out of 5 for keeping the original meaning, 4.23 for natural and professional writing, and 4.38 for usefulness in journalism work. Headline usefulness received 4.54 and ease

of use received 4.38. The full system also received 4.38 for overall usefulness. These ratings show practical value, but they are not model accuracy. The users did not all score the same blind set of outputs. We therefore do not report inter rater agreement. A future blind study should ask at least two or three Sinhala speaking raters to score the same unseen rewrites for style, factual preservation, fluency, and publication usefulness.

E. What the Language Experiments Show

The four tasks reveal different kinds of language failure. Grammar needs restraint as much as correction because a model can damage a clean sentence. Headline generation shows that a good looking score can come from duplicated language examples rather than better generation. Summarization shows that length control and fluency do not guarantee factual safety. Style rewriting shows that a model can learn repeated templates very well if the target data contain them. Across the tasks, these failures have one common implication. Evaluation must account for training data quality and evaluation design.

These results support task specific evaluation instead of one overall metric. A summary can have high word overlap and still change a country. A style output can look different from the source and still lose an important fact. A grammar model can correct more faulty sentences while also editing clean text. For this reason, we report task specific measures and keep human feedback separate from automatic scores.

V. SYSTEM INTEGRATION

The four adapters are connected to one editorial platform through a FastAPI backend and the shared SinLlama server. The web application, browser client, and Google Docs client use the same service. Each task keeps its own output checks. These checks are reported separately from adapter scores so that post processing is not counted as model quality.

VI. MAIN FINDINGS, REPRODUCIBILITY, AND LIMITATIONS

Across all tasks, scores were sensitive to target quality, split construction, and evaluation design. Grammar improved after target repair even though the number of training inputs did not increase. Headline scores were inflated by a split that allowed leakage. Summarization needed article level splitting because three summaries came from the same source article. Style training showed that a model can learn copied and repeated targets very well while still failing the real rewriting task. These findings show why data checks should happen before a higher score is treated as real progress.

Several limitations remain. We trained all adapters with one fixed seed. The grammar confidence intervals show test set uncertainty and not training variation. Our field study gives human feedback but does not replace a blind multi rater style test. Summarization still needs a stronger check for entity and event factuality. The final headline model needs a new evaluation with a fixed generation seed and the exact production prompt. One member of our team checked Stage 6,

so independent language review would make the benchmark stronger.

We do not report one combined accuracy because the four tasks measure different things. Grammar needs exact correction and clean preservation. Headline and summarization use overlap and length measures, while style also needs fact preservation and human feedback. ROUGE and BERTScore are useful signals, but neither proves that every name, number, and event is correct. Keeping the measures separate avoids giving one score more meaning than it has.

We collected the corpus from publicly available news pages for academic research. We kept source and URL information and removed exact duplicates. We will not release full article text when copyright or publisher rules do not allow it. We plan to release filtering and evaluation code, task prompts, adapter settings, split identifiers and hashes, source links or stable hashes, filtered metadata, and result manifests where permitted. We also plan to keep the adapter version, prompt version, decoding settings, seed information, and output hashes with saved evaluation records. These records are important for the headline task because two saved sampled runs later gave different results.

VII. CONCLUSION

We study data centric adaptation of one Sinhala language model for four newsroom tasks. Our grammar results show that task adaptation is needed. They also show why exact correction on faulty text and clean preservation should be reported with the do nothing baseline. Our Stage 6 replay shows that the current safety validator keeps the neural result instead of creating a hidden accuracy gain. Our headline experiment shows how copying rows before splitting can more than double the apparent ROUGE score. Our summarization audit shows why multiple targets from one article must be split at article level before sample expansion. After retraining, the raw and cleaned recipes give very similar results on 273 held out articles. For style rewriting, we add a 500 output frozen automatic test and a small human field study.

Across all four tasks, more data or a newer adapter does not always mean real progress. Target quality, data splits, prompts, and evaluation choices can change the conclusion. The experiments also show Sinhala specific risks. A common repetition control damaged an opening grapheme, while fluent

summaries could still change a country or a number scale. These checks are part of the research result and not only implementation details.

REFERENCES

- [1] N. de Silva, “Survey on Publicly Available Sinhala Natural Language Processing Tools and Research,” arXiv:1906.02358, 2019.
- [2] R. Jayakody and G. Dias, “Performance of Recent Large Language Models for a Low-Resourced Language,” in *Proc. Int. Conf. Asian Language Processing (IALP)*, 2024, pp. 162–167, doi: 10.1109/IALP63756.2024.10661169.
- [3] V. Dhananjaya, P. Demotte, S. Ranathunga, and S. Jayasena, “BERTifying Sinhala - A Comprehensive Analysis of Pre-trained Language Models for Sinhala Text Classification,” in *Proc. LREC*, 2022, pp. 7377–7385.
- [4] T. Ranasinghe *et al.*, “Sinhala Encoder-only Language Models and Evaluation,” in *Proc. 63rd ACL*, 2025, pp. 8623–8636, doi: 10.18653/v1/2025.acl-long.422.
- [5] A. Grattafiori *et al.*, “The Llama 3 Herd of Models,” arXiv:2407.21783, 2024.
- [6] H. W. K. Aravinda *et al.*, “SinLlama - A Large Language Model for Sinhala,” in *Proc. MERCon*, 2025, pp. 617–622, doi: 10.1109/MER-Con67903.2025.11217094.
- [7] H. Hettiarachchi, D. Premasiri, L. R. C. Uyangodage, and T. Ranasinghe, “NSina: A News Corpus for Sinhala,” in *Proc. LREC-COLING*, 2024, pp. 12307–12312.
- [8] M. A. C. Akmal Jahan and K. K. C. Wijesekara, “Automated Text Summarization of Sinhala Online Articles,” *Journal of Science-FAS-SEUSL*, vol. 4, no. 1, pp. 1–15, 2023.
- [9] M. Goonawardena *et al.*, “Automated Spelling Checker And Grammatical Error Detection And Correction Model for Sinhala Language,” in *Proc. SCSE*, 2022, pp. 184–189, doi: 10.1109/SCSE56529.2022.9905126.
- [10] E. J. Hu *et al.*, “LoRA: Low-Rank Adaptation of Large Language Models,” in *Proc. ICLR*, 2022.
- [11] T. Dettmers, A. Pagnoni, A. Holtzman, and L. Zettlemoyer, “QLoRA: Efficient Finetuning of Quantized LLMs,” in *Advances in Neural Information Processing Systems*, vol. 36, pp. 10088–10115, 2023.
- [12] L. Xue *et al.*, “mT5: A Massively Multilingual Pre-trained Text-to-Text Transformer,” in *Proc. NAACL-HLT*, 2021, pp. 483–498.
- [13] T. Hasan *et al.*, “XL-Sum: Large-Scale Multilingual Abstractive Summarization for 44 Languages,” in *Findings of ACL-IJCNLP*, 2021, pp. 4693–4703.
- [14] L. Xue *et al.*, “ByT5: Towards a Token-Free Future with Pre-trained Byte-to-Byte Models,” *Transactions of the Association for Computational Linguistics*, vol. 10, pp. 291–306, 2022.
- [15] C.-Y. Lin, “ROUGE: A Package for Automatic Evaluation of Summaries,” in *Text Summarization Branches Out*, 2004, pp. 74–81.
- [16] T. Zhang, V. Kishore, F. Wu, K. Q. Weinberger, and Y. Artzi, “BERTScore: Evaluating Text Generation with BERT,” in *Proc. ICLR*, 2020.
- [17] P. Ostheimer, M. Nagda, M. Kloft, and S. Fellenz, “Text Style Transfer Evaluation Using Large Language Models,” in *Proc. LREC-COLING*, 2024, pp. 15802–15822.